

SOFT620020.02
Advanced Software
Engineering

Bihuan Chen, Pre-Tenure Assoc. Prof.

bhchen@fudan.edu.cn

<https://chenbihuan.github.io>

Course Outline

Date	Topic	Date	Topic
Sep. 10	Introduction	Nov. 05	Compiler Testing
Sep. 17	Testing Overview	Nov. 12	Mobile Testing
Sep. 24	Holiday	Nov. 19	Delta Debugging
Oct. 01	Holiday	Nov. 26	Presentation 1
Oct. 08	Guided Random Testing	Dec. 03	Bug Localization
Oct. 15	Search-Based Testing	Dec. 10	Automatic Repair
Oct. 22	Performance Analysis	Dec. 17	Symbolic Execution
Oct. 29	Security Testing	Dec. 24	Presentation 2

Discussion – Do You Trust Compilers?



```
int main () {  
    // do something  
    ...  
    return 0;  
}
```



```
...  
call puts  
movl $0, %eax  
popq %rbp  
ret
```

Developers' **belief**: Compilers are **faithful** translators

Reflect on This Trust – llvm bug 14972

```
struct tiny { char c; char d; char e; };

void foo(struct tiny x) {
    if (x.c != 1) abort();
    if (x.e != 1) abort();
}

int main() {
    struct tiny s;
    s.c = 1; s.d = 1; s.e = 1;    removed by optimization
    foo(s);
    return 0;
}
```

```
$ clang -m32 -O0 test.c ; ./a.out
$ clang -m32 -O1 test.c ; ./a.out
Aborted (core dumped)
```

Reflect on This Trust – gcc bug 58731

```
int a, b, c, d, e;
int main() {
    for (b = 4; b > -30; b--)
        int f = 2147483647 - b; integer overflow
        for (; c++)
            for (;;) {
                e = a > f;
                if (d) break;
            }
    return 0;
}
```

```
$ gcc -O0 test.c ; ./a.out
$ gcc -O3 test.c ; ./a.out
^C
```

Reflect on This Trust – Debatable Bug 1

```
char *buf = ...;
char *buf_end = ...;
unsigned int len = ...;

if (buf + len >= buf_end)
    return; // len too large


if (buf + len < buf) → if (0)
return; // overflow, buf + len wrapped around

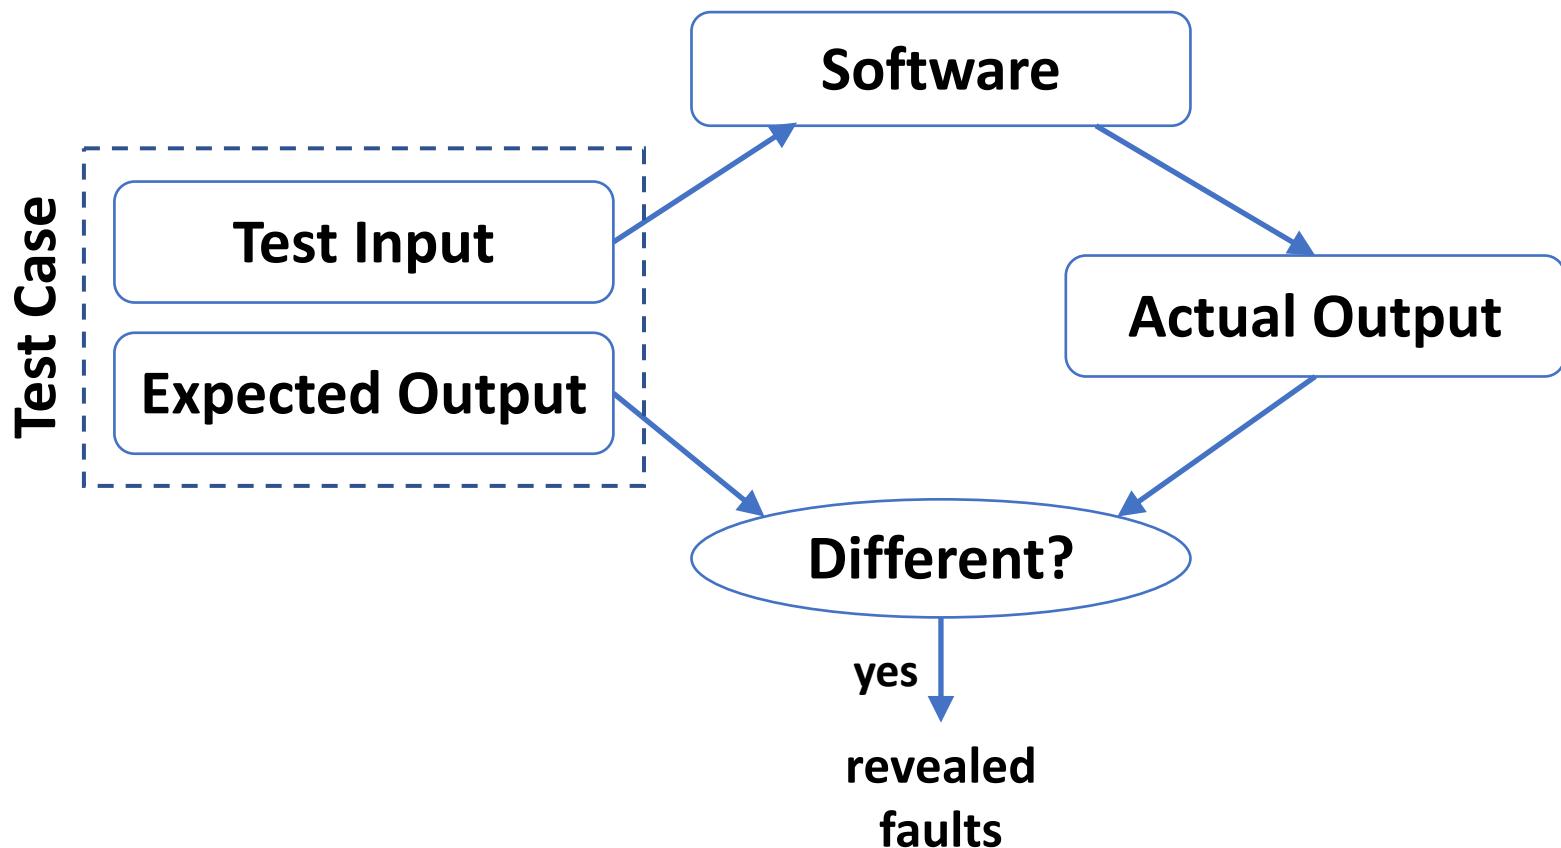
// write to buf[0..len-1]
...
```

But this is a **security threat**, i.e., use a large **len** to trigger **buffer overflow**

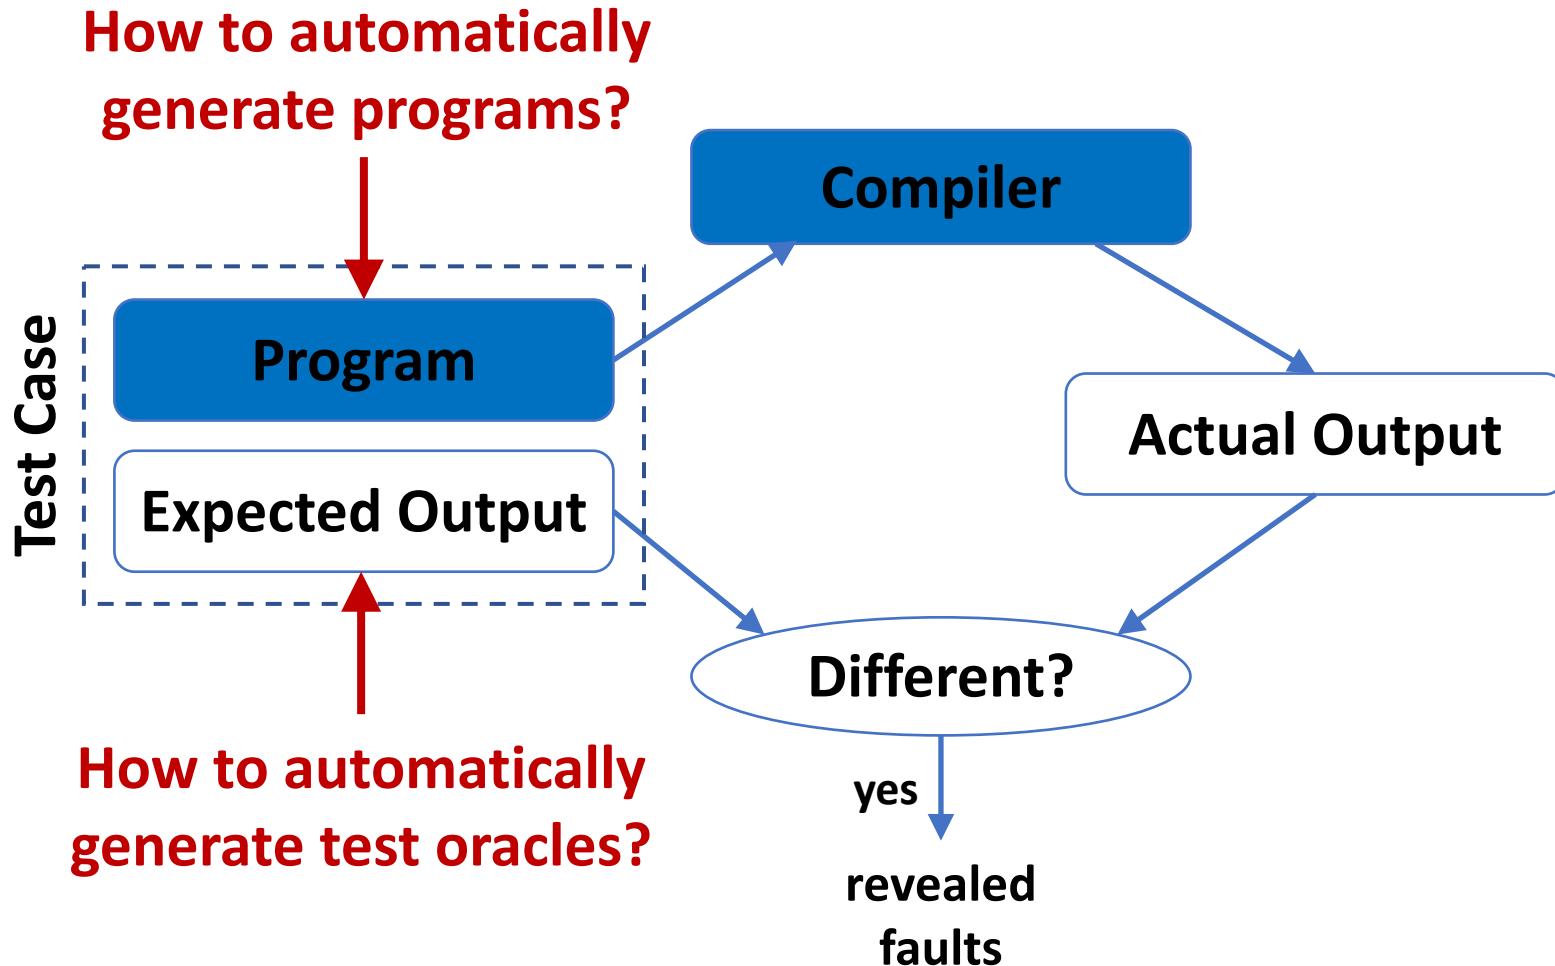
Reflect on This Trust – Debatable Bug 2

```
bool isPasswordCorrect() {  
    bool isPasswordCorrect = false;  
    string Password("password");  
  
    if(Password == getPasswordFromUser()) {  
        isPasswordCorrect = true;  
    }  
    // removed from the optimized code  
    // secure the code by wiping the password from memory  
    memset(Password, 0, sizeof(Password));  
  
    return isPasswordCorrect;  
}
```

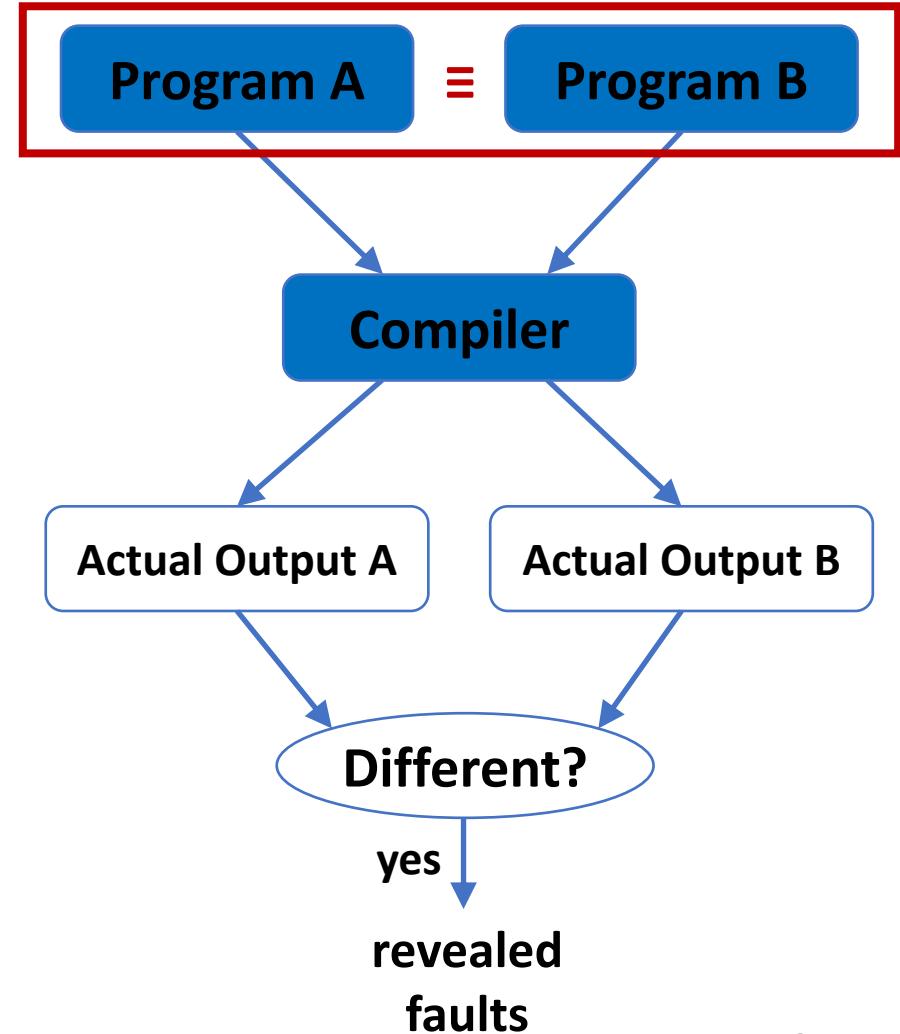
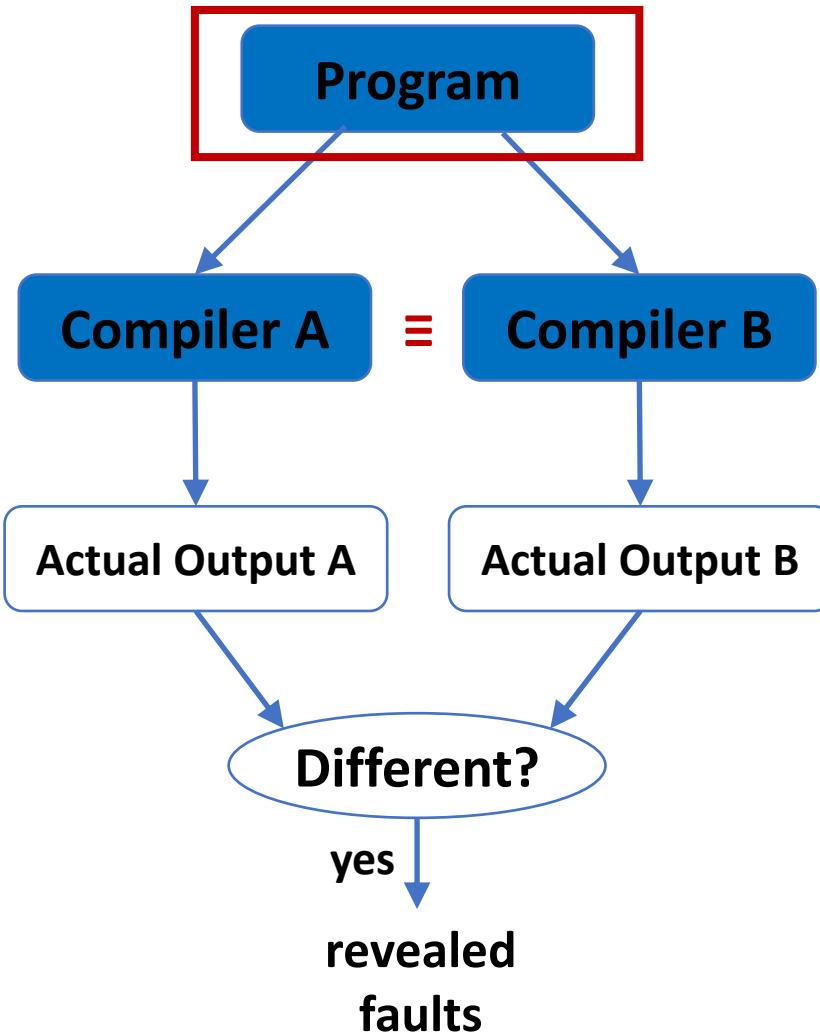
Typical Software Testing



Compiler Testing



Differential Testing (Solving Oracle Problem)

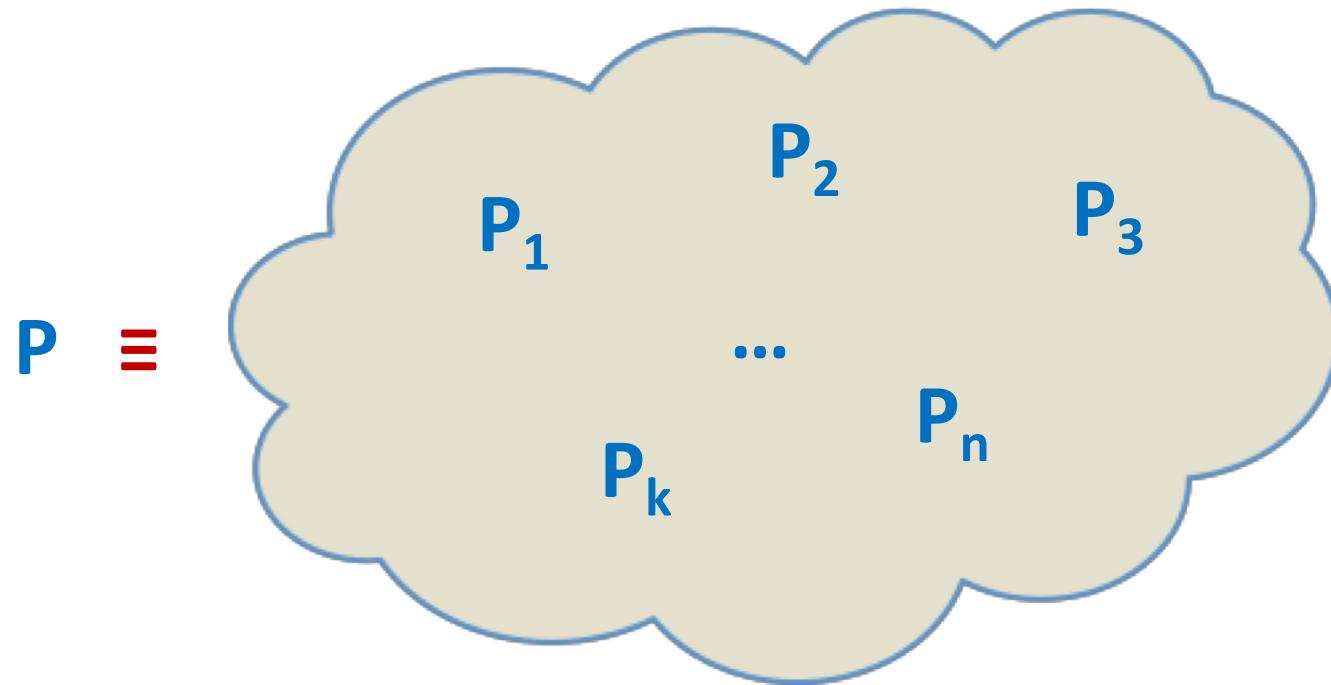


Compiler Validation via Equivalence Modulo Inputs

Vu Le, Mehrdad Afshari, Zhendong Su

PLDI 2014, Distinguished Paper Award

Vision for Equivalent Programs

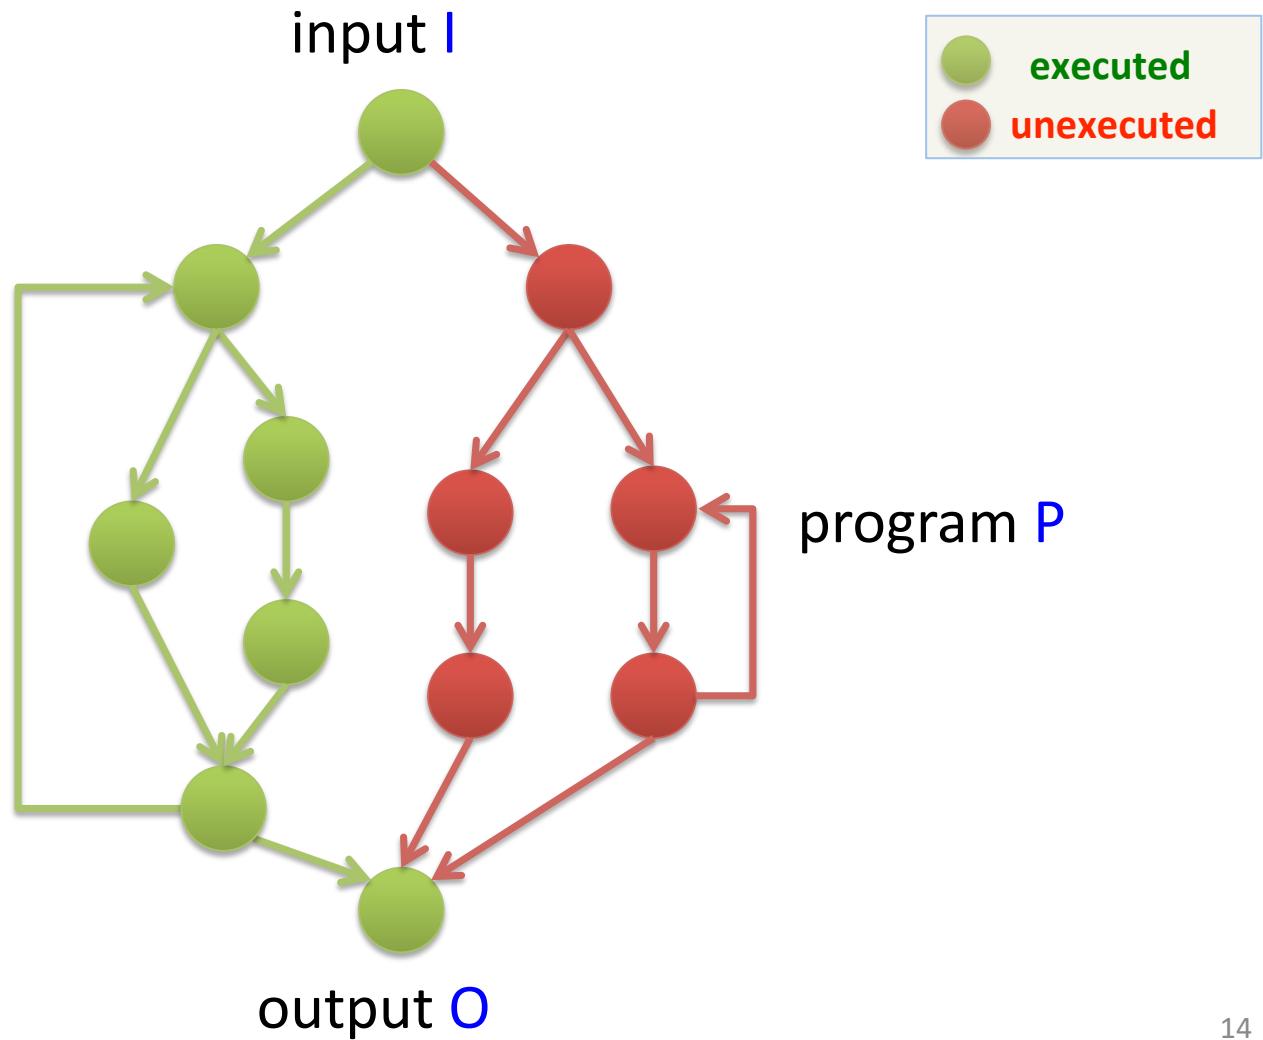


Take existing real-world code and transform it into a different but equivalent variants of the original code

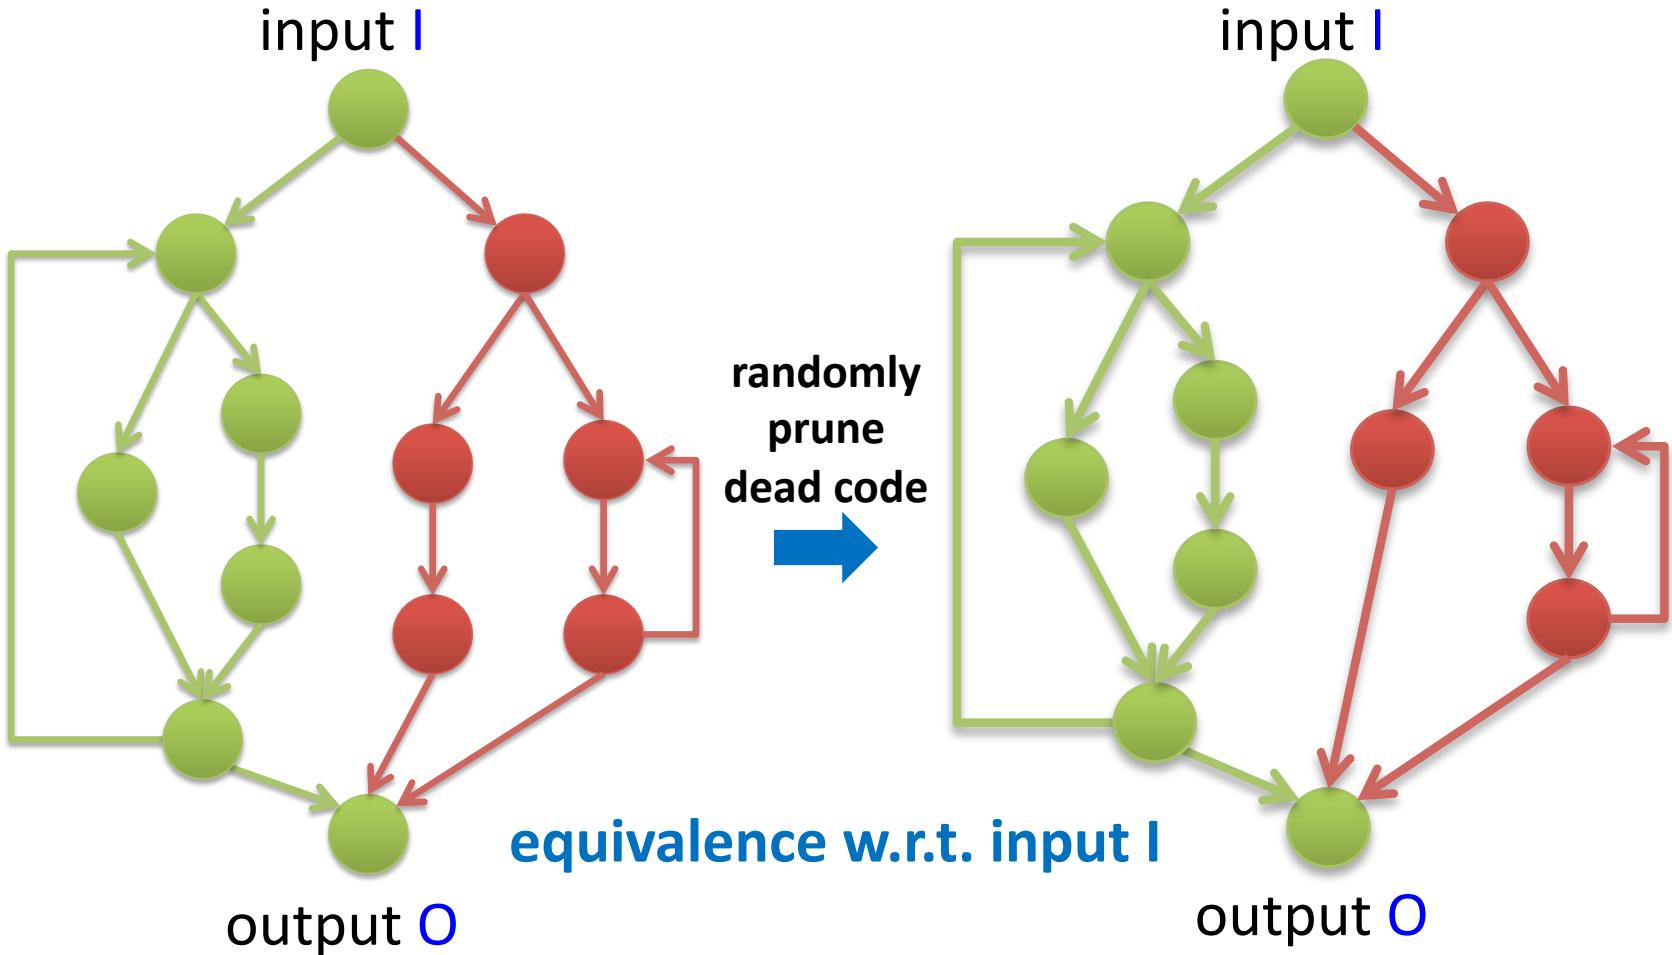
Equivalence Modulo Inputs (EMI)

- Relax equivalence w.r.t. a **given input**
 - Variants must satisfy $P(i) = P_k(i)$ on input i
 - Variants may differ on other input j : $P(j) \neq P_k(j)$
- Exploit the **interplay** between
 - **Dynamic** program execution on **some input**
 - **Static** compilation for **all inputs**

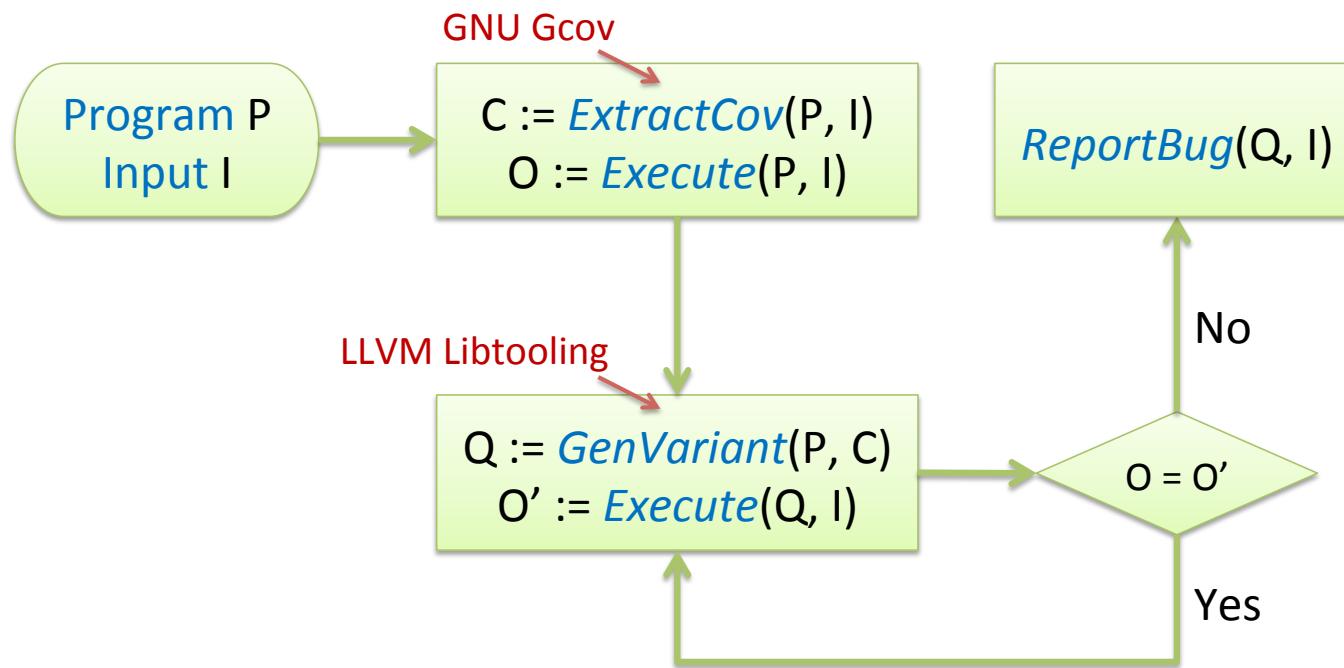
Profile



Mutation



Orion Overview



Examples

```
struct tiny { char c; char d; char e; };
f(int n, struct tiny x, struct tiny y,
  struct tiny z, long l) {
    if (x.c != 10) abort();
    if (x.d != 20) abort();
    if (x.e != 30) abort();
    if (y.c != 11) abort();
    if (y.d != 21) abort();
    if (y.e != 31) abort();
    if (z.c != 12) abort();
    if (z.d != 22) abort();
    if (z.e != 32) abort();
    if (l != 123) abort();
}
main() {
    struct tiny x[3];
    x[0].c = 10;
    x[1].c = 11;
    x[2].c = 12;
    x[0].d = 20;
    x[1].d = 21;
    x[2].d = 22;
    x[0].e = 30;
    x[1].e = 31;
    x[2].e = 32;
    f(3, x[0], x[1], x[2], (long)123);
    exit(0);
}
```

unexecuted

```
struct tiny { char c; char d; char e; };
f(int n, struct tiny x, struct tiny y,
  struct tiny z, long l) {
    if (x.c != 10) /* deleted */;
    if (x.d != 20) abort();
    if (x.e != 30) /* deleted */;
    if (y.c != 11) abort();
    if (y.d != 21) abort();
    if (y.e != 31) /* deleted */;
    if (z.c != 12) abort();
    if (z.d != 22) /* deleted */;
    if (z.e != 32) abort();
    if (l != 123) /* deleted */;
}
main() {
    struct tiny x[3];
    x[0].c = 10;
    x[1].c = 11;
    x[2].c = 12;
    x[0].d = 20;
    x[1].d = 21;
    x[2].d = 22;
    x[0].e = 30;
    x[1].e = 31;
    x[2].e = 32;
    f(3, x[0], x[1], x[2], (long)123);
    exit(0);
}
```

GCC test suite

```
clang -m32 -O0 test.c ; ./a.out
clang -m32 -O1 test.c ; ./a.out
```

```
clang -m32 -O0 test.c ; ./a.out
clang -m32 -O1 test.c ; ./a.out
Aborted (core dumped)
```

Examples (cont.)

```
int a, b, c, d, e;
int main() {
    for (b = 4; b > -30; b--)
        for (; c++)
            for (;;) {
                b++;
                e = a > 2147483647 - b;
                if (d) break;
            }
    return 0;          unexecuted
}
```

Generated by CSmith

```
gcc -O0 test.c ; ./a.out
gcc -O3 test.c ; ./a.out
```

```
int a, b, c, d, e;
int main() {
    for (b = 4; b > -30; b--)
        for (; c++)
            for (;;) {
                e = a > 2147483647 - b;
                if (d) break;
            }
    return 0;
}
```

```
gcc -O0 test.c ; ./a.out
gcc -O3 test.c ; ./a.out
^C
```

Why Effective?

- Compilers produce correct code for **all inputs**
- Variants have different data & control flow
 - Exercise **various** optimization strategies
 - Demand exact **same output** on the given input

Evaluation

- Two multi-core Ubuntu machines
- April 2013 – March 2014
- Seed programs
 - Compiler test suites
 - Open-source projects
 - CSmith-generated programs

Bug Counts and Types

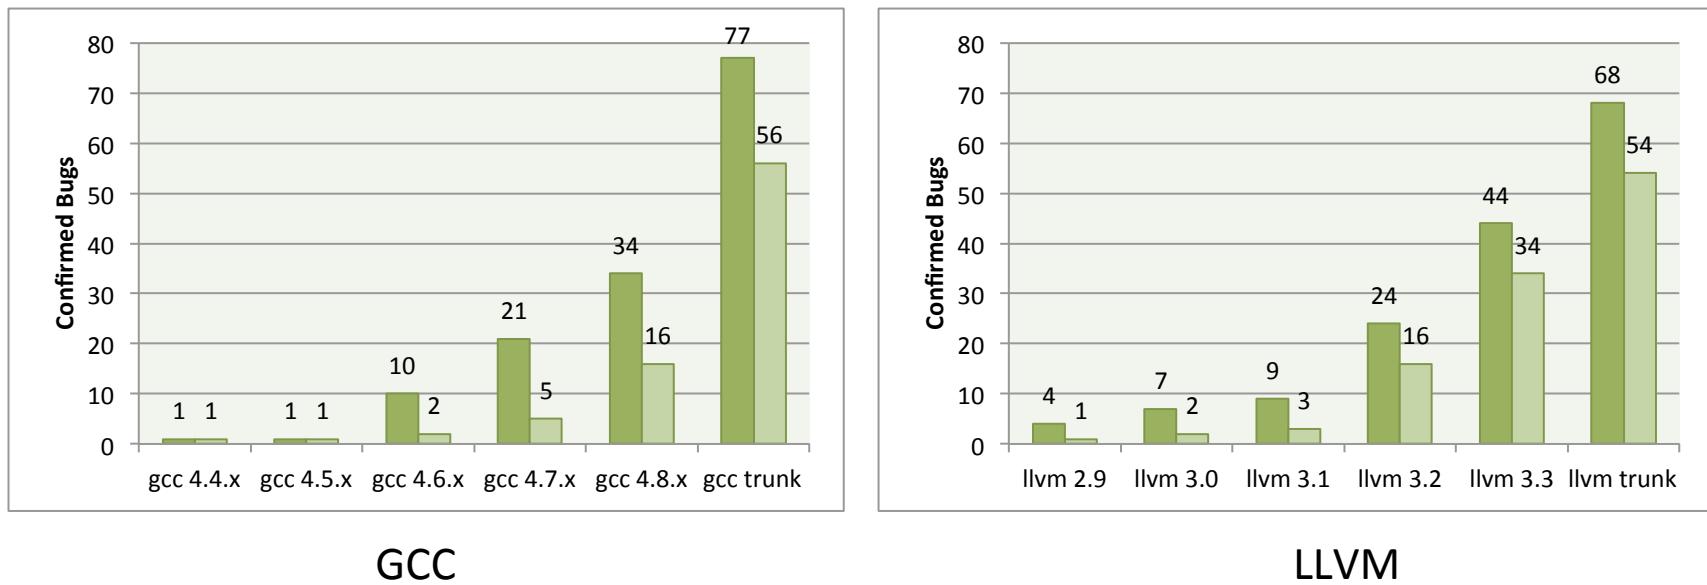
Bug Counts

	GCC	LLVM	TOTAL
Reported	111	84	195
Marked Duplicate	28	7	35
Confirmed	79	68	147
Fixed	56	54	110

Bug Types

	GCC	LLVM	TOTAL
Wrong code	46	49	95
Crash	23	10	33
Performance	10	9	19

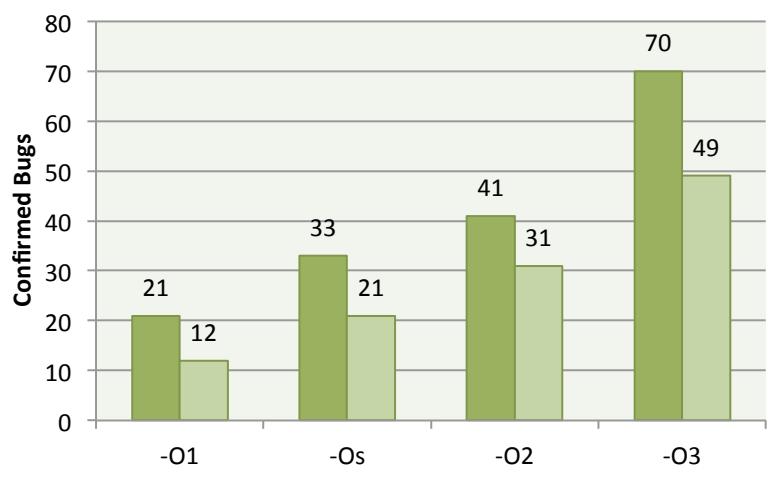
Affected Versions



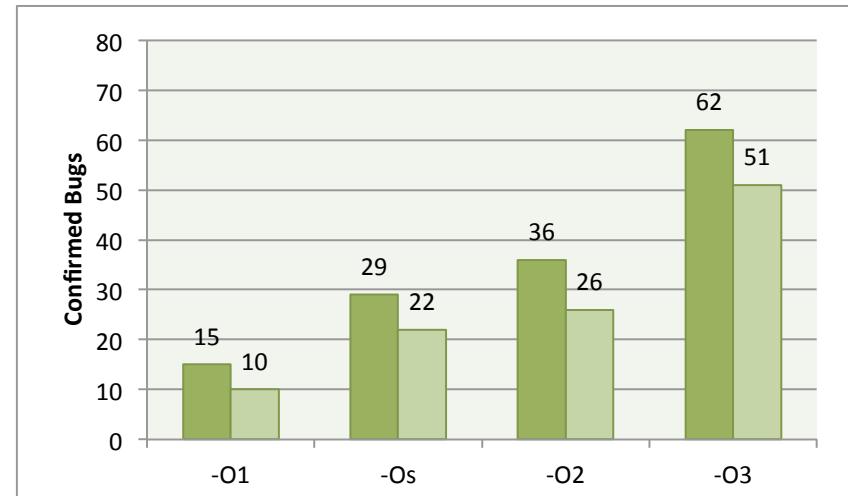
GCC

LLVM

Affected Optimization Levels

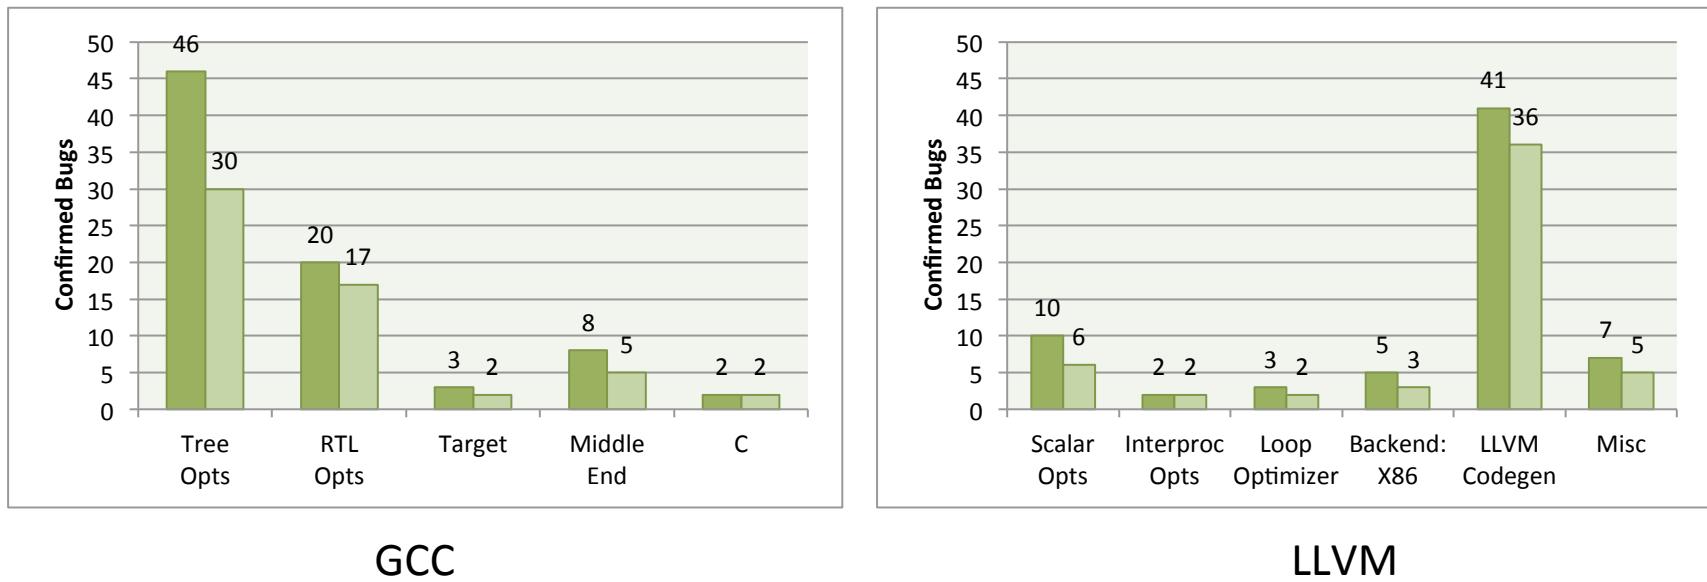


GCC



LLVM

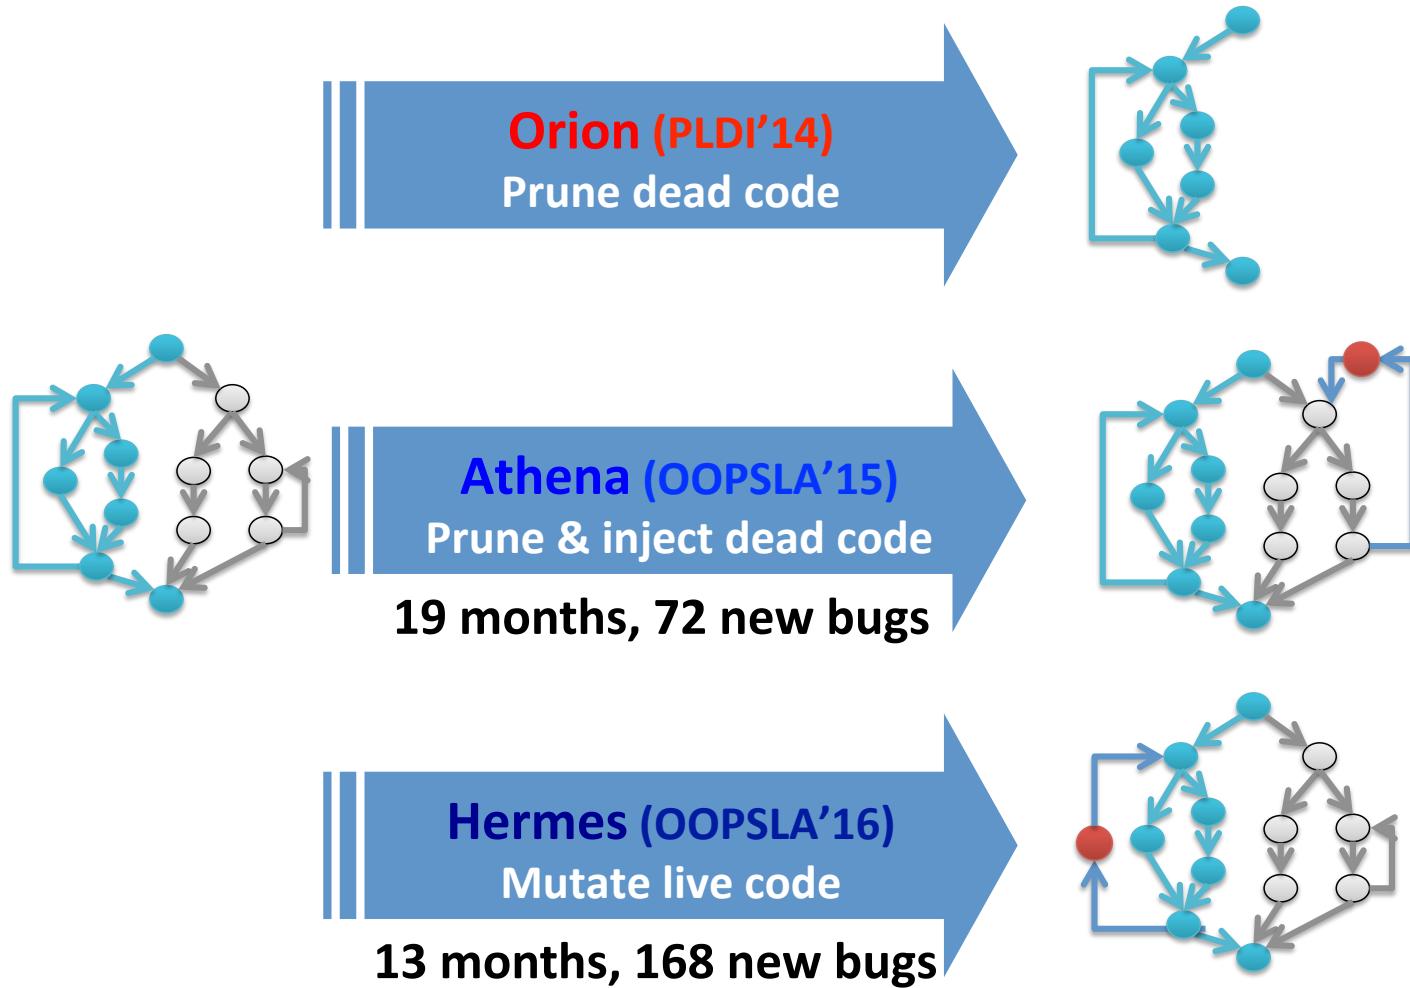
Affected Components



Conclusions

- EMI is **general** and **widely applicable**
 - Can test compilers, analysis and transformation tools
 - Generates real-world tests
 - Requires no reference compilers
- Orion is very **effective**
 - Has uncovered **200+ bugs** in GCC and LLVM
 - Majority of the bugs were **miscompilations**
- Plan to investigate other **mutation strategies**

More Mutation Strategies



Bug Counts (till Recently)

	GCC	LLVM	TOTAL
Reported	643	498	1141
Fixed	397	239	636

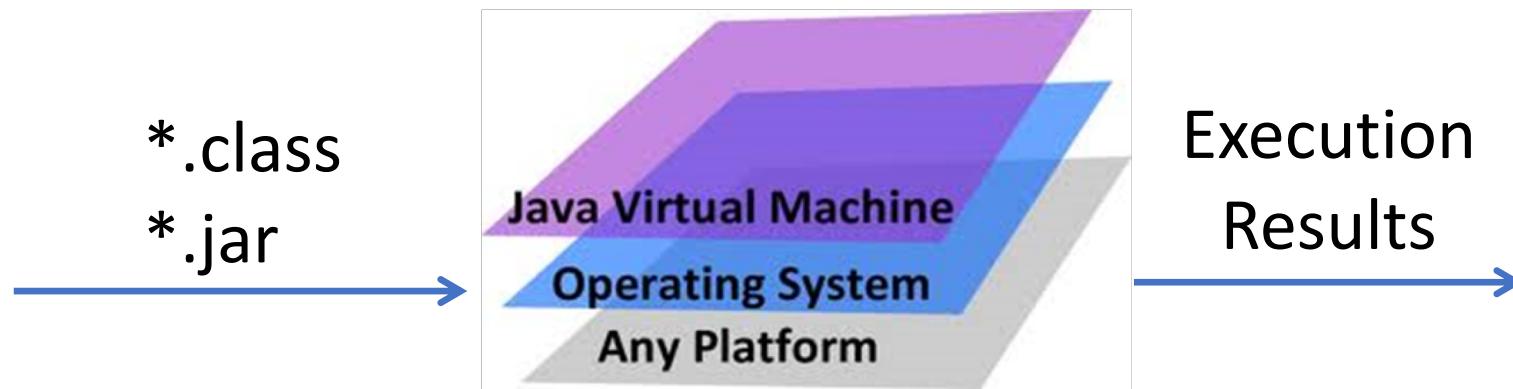
Coverage-Directed Differential Testing of JVM Implementations

Yuting Chen, Ting Su, Chengnian Sun, Zhendong Su, Jianjun Zhao

PLDI 2016

Java Virtual Machine (JVM)

Oracle's HotSpot, IBM's J9, Jikes RVM, Azul's Zulu, and GNU's GIJ



Reflect on This Trust – JVM Discrepancy 1

```
1 ...  
2 MD5 checksum 8  
3 ...  
4 class M1436188543  
5 minor version: 0  
6 major version: 51  
7 flags: ACC_SUPER  
8 Constant pool:  
9 ...  
10 #7 = Utf8 <clinit>  
11 #8 = Utf8 ()V  
12 #9 = Class #19 // java/lang/System  
13 #10 = Utf8 Code  
14 #11 = Utf8 main  
15 ...  
16 {  
17     public abstract {};  
18     flags: ACC_PUBLIC, ACC_ABSTRACT  
19     public static void main(java.lang.String[]);  
20     flags: ACC_PUBLIC, ACC_STATIC  
21     Code:  
22         stack=2, locals=1, args_size=1  
23         0: getstatic #12 // Field java/lang/System.  
             out:Ljava/io/PrintStream;  
24         3: ldc #4 // String Completed!  
25         5: invokevirtual #21 // Method java/io/  
             PrintStream.println:(Ljava/lang/String  
             ;)V  
26         8: return}
```

clinit method

public abstract {};

- HotSpot takes it as an **ordinary** method
- J9 reports a **class format error**

Cause: the JVM specification says that “other methods named <clinit> in a class file are **of no consequence**”



A class or interface initialization method needs to be strictly defined

Reflect on This Trust – JVM Discrepancy 2

```
1 //The Jimple code of M1433982529
2 public class M1433982529 extends java.lang.
3     Object{
4         protected void internalTransform(java.lang.
5             String) {
6             java.util.Map r0;
7             r0 := @parameter0: java.util.Map;
8             staticinvoke <java.lang.Object: boolean
9                 getBoolean(java.util.Map)>(r0);
10            return;
11        }
12    }
```

GIJ throws a **verification error** because it is unsafe to perform a type casting between `java.lang.String` and `java.util.Map`, while HotSpot does not report any error



JVMs take their own classfile verification and type checking policies

Reflect on This Trust – JVM Discrepancy 3

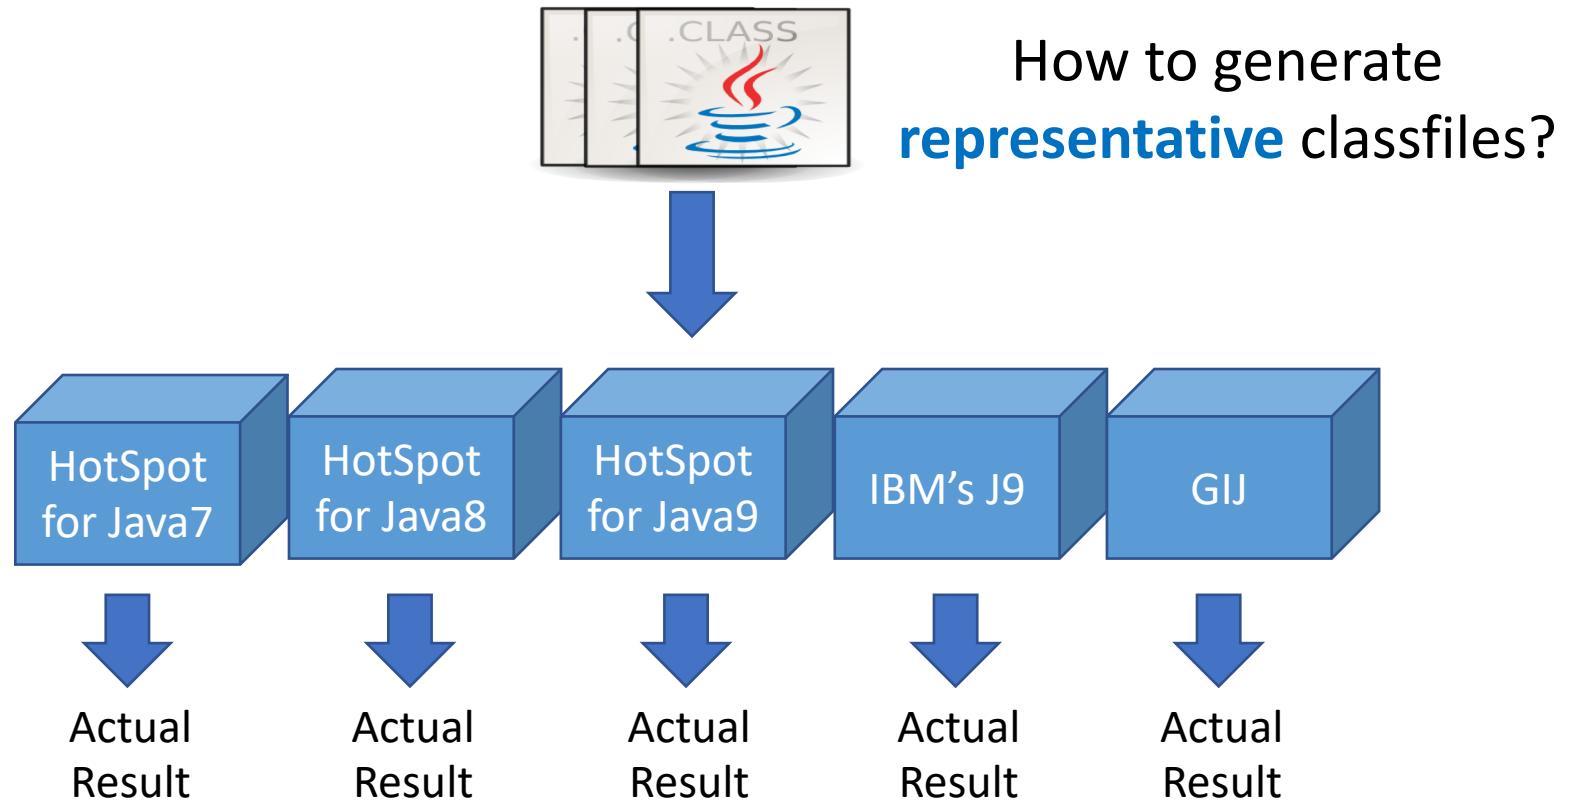
```
1 //The source code of sun.java2d.pisces.  
2 PiscesRenderingEngine  
3 public class PiscesRenderingEngine extends  
4     RenderingEngine {  
5     ...  
6     private static enum NormMode {OFF, ON_NO_AA  
7         , ON_WITH_AA};  
8     ...  
9 }  
//The Jimple code of M1437121261  
10 public class M1437121261 {  
11     public static void main (String[] r0)  
12         throws sun.java2d.pisces.  
13             PiscesRenderingEngine$2{  
14     ...  
15 } }
```

HotSpot reports a `java.lang.IllegalAccessException`, while J9 and GIJ do not



JVMs are not compatible to access some classes

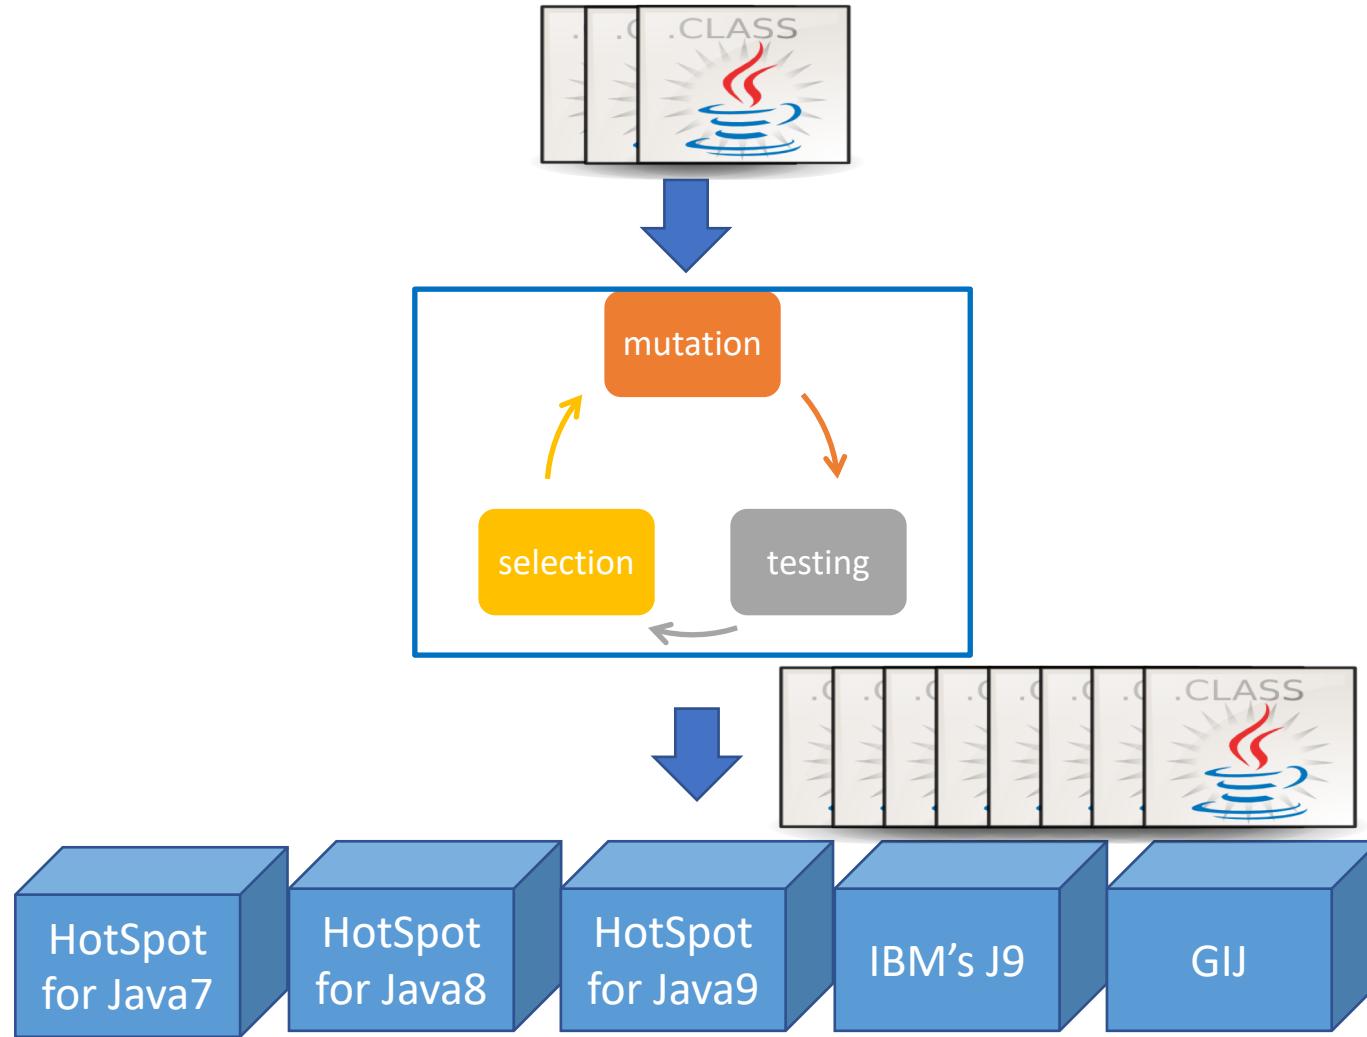
Differential JVM Testing



Two Key Observations

- A classfile can encompass intricate constraints
 - Corner cases can be created through mutating seeds
- Equivalence class partition removes redundancy
 - ECP works only if we can decide whether two tests belong to the same partition

ClassFuzz Overview

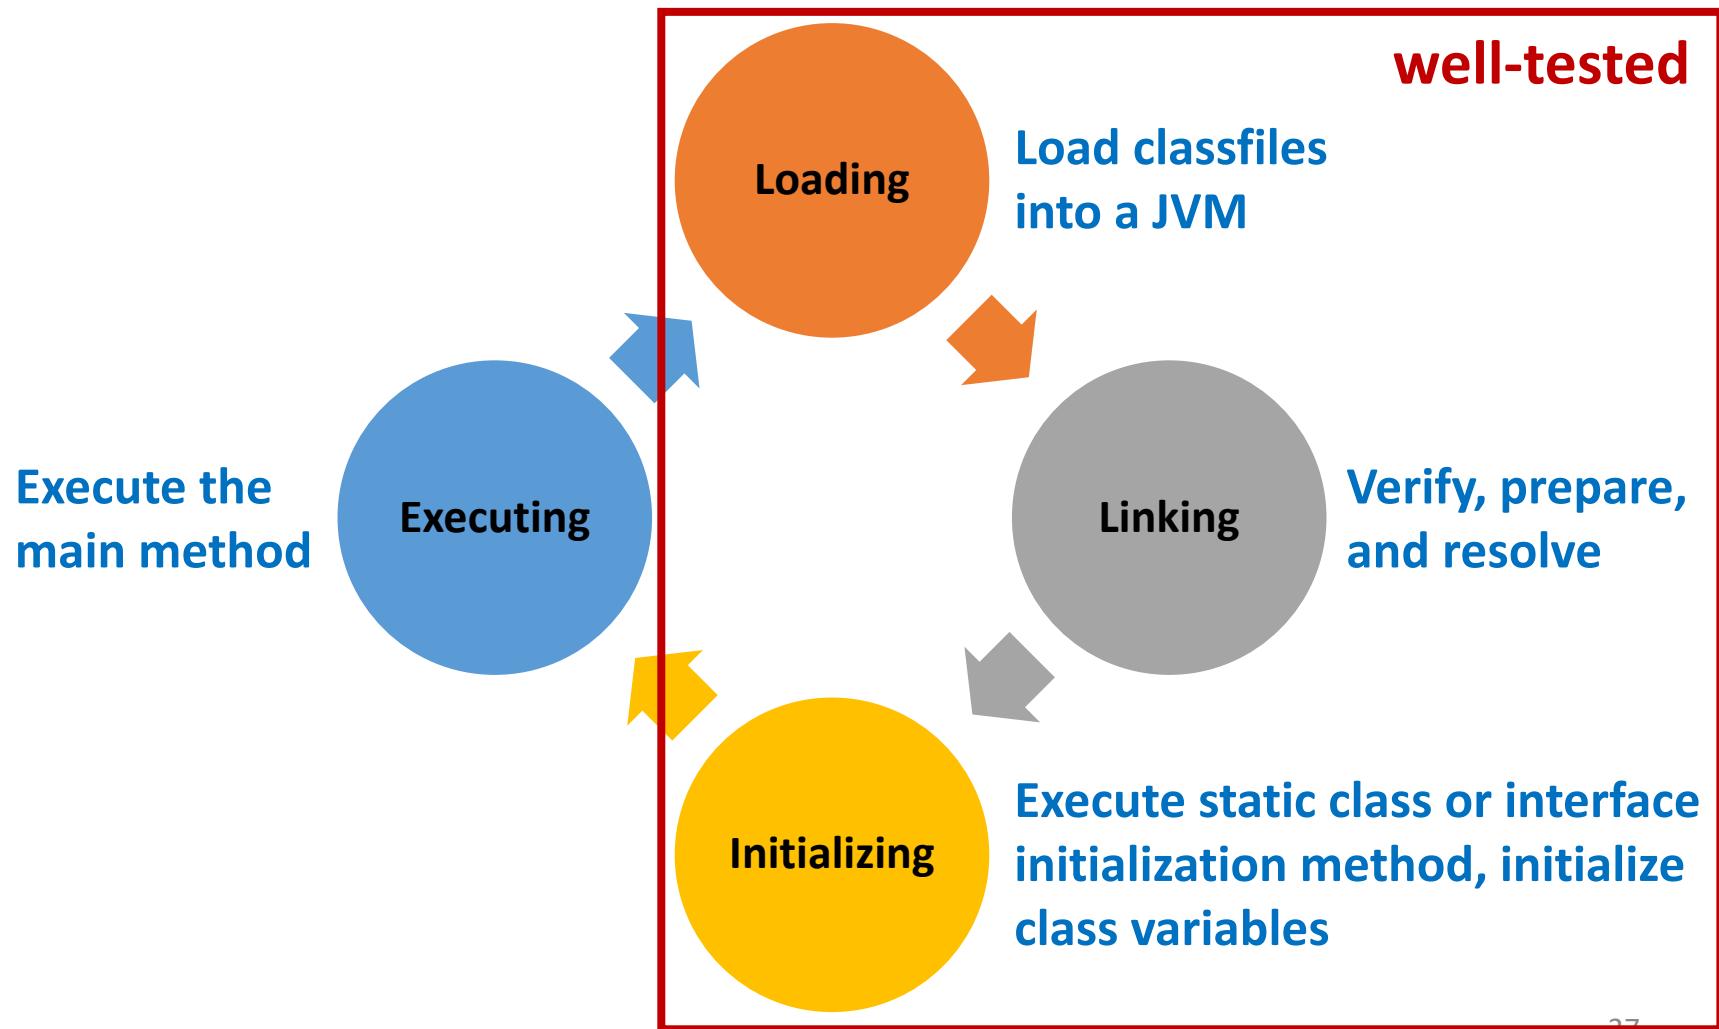


Mutating Classfiles

What to mutate	Typical mutators	Example (Jimple code before mutation → Jimple code after mutation)
Class	Reset its attributes (<i>e.g.</i> , modifiers, name, package name, superclass), <i>etc.</i>	class M1437185190 extends java.lang.Object → private class M1437185190 extends java.lang.Thread
Interface	Insert one or more class-implementing interfaces, delete one or more interfaces, <i>etc.</i>	class M1437185190 extends java.lang.Object → class M1437185190 extends java.lang.Object implements java.security.PrivilegedAction
Field	Insert one or more class fields, delete one or more fields, choose a field and set its attribute(s) (<i>e.g.</i> , name, modifiers), <i>etc.</i>	...protected final java.util.Map MAP; ... → ...protected final java.util.Map MAP; public java.lang.Object MAP; ...
Method	Insert one or more class methods, delete one or more methods, choose a method and set its attribute(s) (<i>e.g.</i> , name, modifiers, return type), <i>etc.</i>	public void <init>(){...} → public static void <init>(){...}
Exception	Select a method and insert one or more exceptions thrown, delete one or more exceptions thrown, <i>etc.</i>	public static void main(java.lang.String[]){ → public static void main(java.lang.String[]) throws sun.java2d.pisces.PiscesRenderingEngine\$2 {
Parameter	Select a method and insert one or more parameters, delete one or more parameters, <i>etc.</i>	public static void main(java.lang.String[]) → public static void main(java.lang.Object , java.lang.String[])
Local variable	Select a method and insert one or more local variables, delete one or more local variables, choose a local variable and set its attribute(s) (<i>e.g.</i> , type, name, modifiers), <i>etc.</i>	public static void main(java.lang.String[]){ int \$i0; ...} → public static void main(java.lang.String[]){ java.lang.String \$i0; ...}
Jimple file	Insert one or more program statements, delete one or more program statements, <i>etc.</i>	r0:=parameter0: java.lang.String[]; r1:=<java.lang.System: java.io.PrintStream out>; virtualinvoke \$r1.<java.io.PrintStream:void println(java.lang.String)>("Executed"); → r0:=parameter0: java.lang.String[]; virtualinvoke \$r1.<java.io.PrintStream:void println(java.lang.String)>("Executed"); r1:=<java.lang.System: java.io.PrintStream out>;

129 mutators

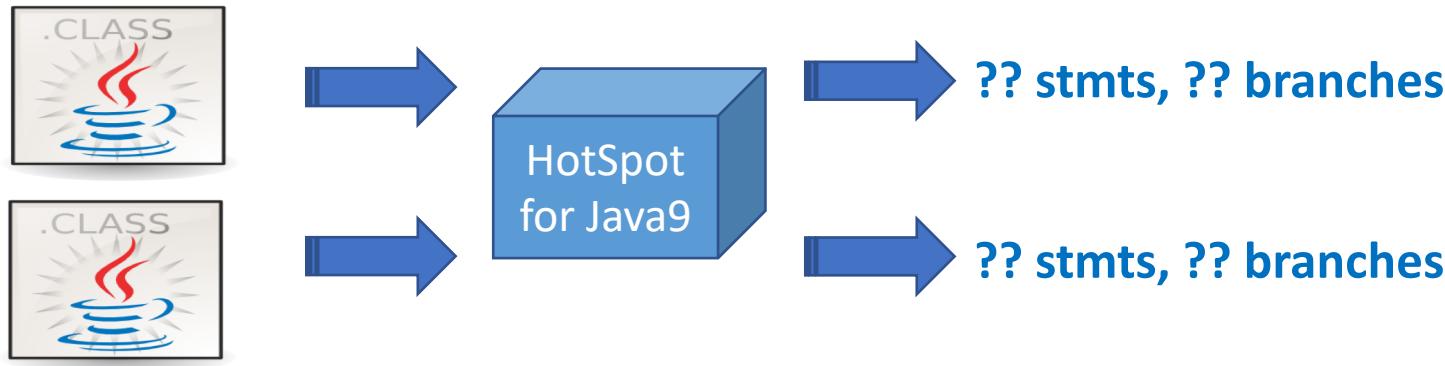
JVM Startup Process



Selecting Mutators

- **Fact:** mutators are designed arbitrarily; and some are effective, while some others are not effective
- **Goal:** create representative classfiles as many and fast as possible
- **Heuristic:** the more number of representative classfiles have been created by a mutator, the more chance should be given for that mutator

Selecting Representative Classfiles



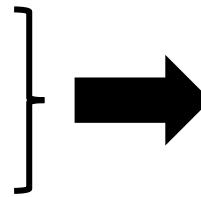
- Several comparison criteria are introduced
 - Statement coverage
 - Branch coverage
 - Trace coverage

Differential Testing of JVMs

- check whether the main method can be normally invoked or during which phase an error/exception is thrown
 - (0) “normally invoked”
 - (1) “rejected during the loading phase”
 - (2) “rejected during the linking phase”
 - (3) “rejected during the initialization phase”
 - (4) “rejected at runtime”

Evaluation Setup

- Coverage collection
 - HotSpot for Java9
 - GCOV + LCOV



At each run the coverage can be conveniently collected

HotSpot (**260K** LOCs)

Cost for cov. analysis: **30+ mins**

share/vm/classfile/ (**11977** LOCs) 

Cost for cov. analysis: **90 secs**

- Seeds
 - 1216 classfiles in JRE 7

Evaluated Setup (cont.)

- ClassFuzz supplemented with a uniqueness analysis
 - [st], [stbr], [tr]
- Randfuzz, Greedyfuzz, Uniquefuzz

	ClassFuzz	randfuzz	greedyfuzz	uniquefuzz
Mutation-based	✓	✓	✓	✓
Cov. Analysis	✓	✗	✓	✓
Uniqueness Analysis	[st] [stbr] [tr]	✗	[stbr]	[stbr]
Mutator Selection	✓	✗	✗	✗

How Many Test Classfiles can be Generated?

	Classfuzz			Uniquefuzz	Greedyfuzz	Randfuzz
	[stbr]	[st]	[tr]			
#iterations	2130	2108	1971	1898	1911	46318
GenClasses	1539	1543	1450	1436	1432	29523
TestClasses	898	494	774	628	98	29523
succ	42.2%	23.4%	39.3%	33.1%	5.1%	63.7%
Average time for each generated class (sec.)	168.4	168.0	178.8	180.5	181.0	8.78
Average time for each test class (sec.)	288.6	524.7	334.9	412.7	2644.9	8.78

$$\text{succ} = |\text{testClasses}| / \#iterations * 100\%$$

How Effective are the Test Classfiles?

	JRE7 classes	seeding classfiles	Classfuzz			Uniquefuzz	Greedyfuzz	Randfuzz
			[stbr]	[st]	[tr]			
GenClasses	21736	1216	1539	1543	1450	1436	1432	
all invoked	77	17	197	199	226	203	195	
all rejected at the same stage	21295	1162	1100	1104	984	1063	980	—
Discrepancies	364	37	242	240	240	170	257	
Distinct_Discrepancies	16	9	17	14	13	10	15	
diff	1.7%	3.0%	15.7%	15.6%	16.5%	11.8%	17.9%	
TestClasses			898	494	774	628	98	29523
all invoked			121	62	122	109	11	3914
all rejected at the same stage			670	374	564	462	72	20072
Discrepancies	—	—	107	58	88	57	15	5537
Distinct_Discrepancies			17	11	12	9	7	14
diff			11.9%	11.7%	11.4%	9.1%	15.3%	18.8%

$$\text{diff} = |\text{Discrepancies}| / |\text{Classes}| * 100\%$$

Can the Test Classfiles Find JVM Defects?

- 62 new JVM discrepancies

```
1 ...  
2 MD5 checksum 8  
3 fb69050bbcb9a83ddd90ae393368c5e  
4 ...  
5 class M1436188543  
6 minor version: 0  
7 major version: 51  
8 flags: ACC_SUPER  
9 Constant pool:  
10 ...  
11 #7 = Utf8 <clinit>  
12 #8 = Utf8 ()V  
13 #9 = Class #19 // java/lang/System  
14 #10 = Utf8 Code  
15 #11 = Utf8 main  
16 rename abstract method as <clinit>  
17 {  
18     public abstract {};  
19     flags: ACC_PUBLIC, ACC_ABSTRACT  
20     public static void main(java.lang.String[]);  
21     flags: ACC_PUBLIC, ACC_STATIC  
22     Code:  
23         stack=2, locals=1, args_size=1  
24         0: getstatic #12 // Field java/lang/System.  
25             out:Ljava/io/PrintStream;  
26         3: ldc #4 // String Completed!  
27         5: invokevirtual #21 // Method java/io/  
28             PrintStream.println:(Ljava/lang/String  
29             ;)V  
30         8: return  
31 }
```

set the parameter type from Map to String

```
1 //The Jimple code of M143982929  
2 public class M143982929 extends java.lang.  
3     Object{  
4         protected void internalTransform(java.lang.  
5             String) {  
6             java.util.Map r0;  
7             r0 := @parameter0: java.util.Map;  
8             staticinvoke <java.lang.Object: boolean  
9                 getBoolean(java.util.Map)>(r0);  
10            return;  
11        } }  
12 }
```

add a throw exception to the main method

```
1 //The source code of sun.java2d.pisces.  
2 PiscesRenderingEngine  
3 public class PiscesRenderingEngine extends  
4     RenderingEngine {  
5     ...  
6     private static enum NormMode {OFF, ON_NO_AA  
7         , ON_WITH_AA};  
8     ...  
9     }  
10    //The Jimple code of M1437121261  
11    public class M1437121261 {  
12        public static void main (String[] r0)  
13            throws sun.java2d.pisces.  
14                PiscesRenderingEngine$2{  
15                    ...  
16                } }  
17 }
```

Conclusions

- **Problem**
 - Testing JVMs requires painstaking effort in designing **test classfiles** along with their **test oracles**
- **Proposal: coverage-directed fuzz testing**
 - **Test classfile generation**
 - Mutating classes and selectively applying mutators
 - Deciding the representativeness of a classfile mutant
 - **Differential JVM testing**

Reading Materials

- V. Le, M. Afshari, and Z. Su. Compiler validation via equivalence modulo inputs. In PLDI, page 216-226, 2014.
- V. Le, C. Sun, and Z. Su. Finding deep compiler bugs via guided stochastic program mutation. In OOPSLA, pages 386– 399, 2015.
- C. Sun, V. Le, and Z. Su. Finding compiler bugs via live code mutation. In OOPSLA, pages 849–863, 2016.
- Q. Zhang, C. Sun, and Z. Su. Skeletal program enumeration for rigorous compiler testing. In PLDI, page 347-361, 2017.
- X. Yang, Y. Chen, E. Eide, and J. Regehr. Finding and understanding bugs in C compilers. In PLDI, pages 283–294, 2011.
- Y. Chen, A. Groce, C. Zhang, W. Wong, X. Fern, E. Eide, and J. Regehr. Taming compiler fuzzers. In PLDI, pages 197–208, 2013.
- C. Sun, V. Le, and Z. Su. Finding and analyzing compiler warning defects. In ICSE, pages 203–213, 2016.
- Y. Chen, T. Su, C. Sun, Z. Su, and J. Zhao. Coverage-directed differential testing of JVM implementations. In PLDI, page 85-99, 2016.

Q&A?

Bihuan Chen, Pre-Tenure Assoc. Prof.

bhchen@fudan.edu.cn

<https://chenbihuan.github.io>